

# GestureSpace Project Analysis

## Creative Background

GestureSpace is the result of unifying three previously separate research projects into a single assistive technology solution. The first component was an AprilTag-based object localization system developed to determine camera position and orientation using fiducial markers <sup>1</sup> <sup>2</sup>. This provided a way to map real-world coordinates by placing known AprilTag markers in the environment and using computer vision to locate the camera relative to them. The second component stemmed from a gesture-based computer control project, where hand motions (like pointing or specific hand signs) were tracked using MediaPipe and interpreted by a custom neural network. In that project, a convolutional model was trained on MediaPipe hand landmark data to classify gestures <sup>3</sup> <sup>4</sup>, enabling control of computer functions via simple hand movements. The third influence was the **Aurora Echo** feedback concept – essentially a screen-reader style audio feedback system. This provided the motivation for delivering information to the user through spoken descriptions, similar to how traditional assistive tools read out screen text for visually impaired users. In GestureSpace, this idea manifests as a text-to-speech pipeline that voices environmental information to the user <sup>5</sup>. By integrating these three threads – spatial localization, gesture recognition, and audio feedback – GestureSpace was conceived as a comprehensive assistive solution. The origin story reflects a creative synthesis: experiences from an AI summer camp (where MediaPipe and neural nets were explored) and a winter innovation camp (focusing on camera pose estimation with AprilTags) were combined <sup>6</sup>. The motivation was to help users, especially those who are visually impaired, interact with and understand their physical surroundings in a natural way. GestureSpace's development was driven by the question: *What if we could point at something, ask about it, and hear a helpful answer?* This vision led directly to the current solution that merges voice, vision, and touchless interaction into one system.

## Principle of Work

GestureSpace operates through a multi-stage pipeline that integrates voice recognition, computer vision, spatial calculations, and language generation. The interaction begins with **wake word detection**. A small offline engine (Porcupine by Picovoice) continuously monitors the microphone for a specific wake phrase <sup>7</sup>. When the user speaks the chosen keyword (for example, “Gesture Space”), the Porcupine engine detects it and triggers the system to start active processing. This wake word mechanism allows the system to remain dormant until needed, conserving resources and ensuring hands-free activation.

Once “awake,” the system enters the **speech recognition** stage. It uses a microphone (via PyAudio) to record the user’s voice query, automatically stopping when the user finishes speaking. A noise threshold and timing mechanism are implemented to detect a period of silence indicating the end of the question <sup>8</sup>. The recorded audio is then transcribed by a pretrained neural network model for Automatic Speech Recognition (ASR). Specifically, GestureSpace employs Facebook’s Wav2Vec 2.0 model (the large 960h version) to convert spoken words into text <sup>9</sup>. This deep learning model is known for its accuracy in understanding spoken language and operates here to produce the text of the user’s request (for example, the user might ask, “What objects are on the table to my left?”). Using an advanced ASR model allows the

system to handle natural, unrehearsed questions from the user. The output of this stage is a text string of the user's question or command, now ready for the vision module to address.

In the **visual analysis** stage, GestureSpace captures an image of the user's surroundings through a camera (such as a webcam or smart glasses camera). Two parallel vision processes then occur on this image. First, an object detection model identifies and locates general items in view. GestureSpace uses a YOLO-based convolutional neural network to perform real-time object detection. (In the repository, a model weight file named "yolo12l.pt" is loaded, indicating a YOLO vX model variant is used <sup>10</sup>.) YOLO – which stands for "You Only Look Once" – is an algorithm family known for fast and accurate detection of multiple objects in a single pass through the neural network. The model processes the image and outputs bounding boxes, class names, and confidence scores for each detected object <sup>11</sup>. For example, it might detect a "cup" or "chair" in the scene and provide the coordinates of each in the image. Simultaneously, GestureSpace runs a **MediaPipe** hand-tracking solution to interpret any user hand gestures visible to the camera. MediaPipe's pipeline locates 21 keypoints on the user's hand in real time. From these keypoints, GestureSpace focuses on the index finger's orientation: it computes the 3D direction in which the user is pointing <sup>12</sup>. Essentially, the system can tell where the user's finger is directed in the real world. This allows the assistant to know which area or object the user might be referencing or interested in. The fusion of YOLO object data and MediaPipe gesture data means the system not only sees *what* objects are present, but also infers *where* the user's attention is aimed. This is a major innovation over traditional single-modality systems.

Another critical part of visual analysis is **AprilTag detection for pose estimation**. AprilTags are small barcode-like fiducial markers that can be placed in the environment as reference points. GestureSpace leverages AprilTags to establish a world coordinate frame. When the camera sees these tags, it uses a library (such as `pupil_apriltags`) to detect them in the image <sup>1</sup>. Each tag has a known ID and pre-defined real-world coordinates in a map (for example, the system knows tag ID 0 is at position (x=0.0, y=1.0) in meters in the room's layout) <sup>13</sup>. By detecting multiple AprilTags and knowing their fixed positions, the system can calculate the **camera's exact position and orientation** in the room using solvePnP (Perspective-n-Point solving) from OpenCV <sup>2</sup>. In simpler terms, the camera acts like the user's "eyes," and the AprilTags are like landmarks that help the system figure out where those "eyes" are located in a 3D space. This procedure is akin to how a driver uses known landmarks to understand their location on a map. The result of this step is a transformation matrix or pose (translation and rotation) that converts coordinates between the camera's view and the real-world frame.

With the camera pose known, GestureSpace performs **coordinate translation** for the objects detected. Every object's image coordinates (pixel location in the 2D image) are projected out into the 3D world. The system translates the center of each object's bounding box from image pixels into real-world coordinates (x, y, z) relative to the room and the user <sup>14</sup>. This involves a bit of math: using the camera calibration (intrinsic parameters) to cast a ray from the camera through the pixel into space, and intersecting that ray with a reference plane or using depth cues. In this implementation, a ground plane Z=0 (for example, the floor level) is assumed for intersection <sup>15</sup>. So if an object (say a chair) is detected in the image, the system can estimate that the chair is, for instance, 2.0 meters in front of the user and 1.0 meter to their right, on the floor. All objects' positions can thus be described in relative terms. Moreover, if the user is pointing (from the MediaPipe step), the system can calculate what real-world direction that finger corresponds to and which object lies along that line of pointing. This enables a very intuitive query like pointing at something and asking "What is this?" – the system can match the pointing direction to an object in that area. Traditional screen-reader tools or basic object identifier apps lack this capability; they might tell a user that an object is

present but not *where* it is. In contrast, GestureSpace's use of spatial coordinates lets it convey layout and direction – a key difference that enables spatial reasoning support for the user.

The final stage is **environmental description and feedback**. By now, the system has a list of detected objects with their names and 3D coordinates, and it has the user's spoken question in text form. This information is passed to a large language model (LLM) to generate a helpful response. GestureSpace's design uses a local LLM (in the repo, a model called Qwen 0.6B is loaded) to interpret the scene and answer the user <sup>16</sup>. A prompt is constructed that provides the LLM with the list of objects and their coordinates, and instructs it to act as an assistant describing the scene to a blind person <sup>17</sup>. Notably, the prompt tells the model **not** to read out raw coordinates or numbers, but to use qualitative terms like "near" or "far" when describing distances <sup>17</sup>. This is an important usability feature – a traditional screen reader might read off numbers or GUI labels, but GestureSpace aims to give a natural-language description of physical space (e.g., "There is a chair near the far wall to your left"). The LLM essentially takes the factual data from the vision system and the context of the user's question, and formulates an answer in everyday language. For example, if the user asked, "Is there an open seat here?", the system might respond with something like, "Yes, about two meters in front of you slightly to the right, there is an empty wooden chair." The language model's output is then converted to speech audio using a text-to-speech engine <sup>5</sup> and played through speakers or headphones for the user to hear. The voice feedback closes the loop, giving the user an immediate, spoken answer.

In summary, GestureSpace's principle of work is a **closed-loop interactive system**: voice input → speech-to-text → vision and spatial understanding → text-to-speech output. It can be thought of as a hybrid of a smart voice assistant and an augmented reality system. Unlike traditional screen-readers that only read interface text or simple object detectors that just name objects, GestureSpace performs **multi-modal fusion**. It recognizes what the user said, sees what they see (and what they point at), figures out spatial relationships, and speaks an informative answer. This original combination means the user can ask spatial questions ("Where is...?", "What is in front of me?") and get meaningful answers – something beyond the scope of classical assistive tools. GestureSpace, therefore, offers a richer interaction paradigm for visually impaired users, effectively serving as a talking guide that understands both language and space, whereas conventional tools might only provide one or the other. <sup>17</sup>

## Procedure of Development

The development of GestureSpace proceeded in stages, each integrating a different technology, with iterative testing and refinement at every step. The **first stage** focused on the core vision and localization capabilities. Early on, a simple AprilTag localization demo was created to prove that the camera's pose could be determined reliably. Using a set of AprilTag markers and OpenCV, the team confirmed they could compute the camera's position in real time by detecting multiple tags and solving the PnP problem (mapping 3D points to 2D image points) <sup>2</sup>. They prepared a calibration for the camera (focal lengths, distortion, etc.), stored in a file (`assets/interstices.npz`), to ensure accurate real-world measurements <sup>18</sup>. At this stage, the output was likely just printed coordinates of the camera and perhaps drawn tag outlines on the video feed – a technical foundation demonstrating spatial awareness.

In parallel, a **gesture recognition module** was developed. This involved using Google's MediaPipe to track hand landmarks and designing a neural network to classify gestures. Initial experiments utilized the 21-point hand skeleton from MediaPipe to capture finger positions over time. A custom dataset of hand movements or poses was recorded, and a small CNN (Convolutional Neural Network) was trained to

recognize a set of gesture commands (the repository's `GestureClassifier` suggests four gesture classes as output <sup>19</sup> <sup>20</sup>). Training code was written to augment data (rotations, flips) and extract key points as input features <sup>21</sup> <sup>3</sup>. The training process was iterative: if the accuracy was low, more sample data or network tuning was done. Once the gesture model performed well on test data, its weights were saved to a file (`gesture_classifier.pth`). The code checks for this file at runtime and alerts the developer if the model isn't found, prompting them to train it first <sup>4</sup>. This indicates that a significant development step was collecting training data and refining the gesture classifier until it was reliable for real-time use.

The **object detection** capability was the next piece to integrate. Instead of writing a detector from scratch, the developers opted for the YOLO family of models due to its speed and accuracy in detecting many object types. Integration involved choosing a specific YOLO model and weight configuration. In the code, they instantiate a YOLO model from the Ultralytics library, loading a weights file named `yolo12l.pt` <sup>10</sup>. This suggests they may have used a YOLOv8 Large model or a custom-trained variant (the naming "12l" could imply an experimental version or just an internal naming convention). During development, this step would include verifying that the model can run on the target hardware (possibly needing GPU support for fast inference) and that it correctly identifies objects of interest in the test environment. They likely tested this by printing out detected classes and confidence scores to ensure the model was working as expected <sup>22</sup>. An important challenge was combining the object detector's results with the AprilTag coordinate system. This required careful calibration: the camera intrinsic parameters (which were loaded earlier) needed to be applied so that pixel coordinates from YOLO could be translated to world coordinates. The team had to ensure that the coordinate translation function (`get_point_3d_place`) was accurate – a task that might involve testing with known object positions and adjusting the approach if results were off. The use of a fixed reference plane ( $Z_0 = 0$  for the floor) was decided on to simplify this mapping <sup>23</sup>. Development here likely entailed verifying that when an object was at a known spot (e.g., a bottle placed at a marked location on a table), the system's reported coordinates for that object were reasonably close.

With vision (object + tags) and gesture modules in hand, the development moved to the **voice interface**. Adding voice interaction meant incorporating both input (speech recognition) and output (speech synthesis). The team selected Picovoice Porcupine for wake word detection due to its offline operation and low latency. Setting this up involved obtaining an access key and training or selecting a wake word ("Gesture Space") model file <sup>7</sup>. Development tests here included saying the wake word in various conditions to ensure the system woke up reliably without false positives. For speech-to-text, they integrated the Wav2Vec 2.0 model from HuggingFace. This decision was likely driven by Wav2Vec's strong accuracy without requiring an internet connection. The integration required installing the model and its tokenizer, and ensuring audio input from the microphone could be fed into it correctly. They wrote a recording loop with silence detection to segment the user's speech <sup>8</sup>, which they tested by speaking sample queries. For instance, they might say "What's on my desk?" and then check if the transcription text was accurate. Challenges here could include managing audio buffering and avoiding clipping the start or end of the user's speech. Once the ASR was returning correct text, they connected it to the rest of the pipeline – i.e., using the recognized question to decide what visual information to gather.

The **language model and response generation** came last in development. Initially, the system may have provided very raw responses (like simply reading out object names and distances). To improve the user experience, the team integrated a large language model to generate more natural descriptions. They experimented with a small local model (Qwen-3 with 0.6B parameters) for quick iteration <sup>24</sup>. They crafted a system prompt (found in `shared.py`) instructing the model to describe the scene to a blind person in

qualitative terms <sup>17</sup>. Development involved testing various prompt phrasings and ensuring the model's output was relevant. For example, if the user asked "Where are the exits?", the prompt needed to guide the model to use the content (maybe it detected doors in the object list) and respond accordingly. They likely also implemented a fallback: the code shows a mechanism where if the local model returns a special token `<remote>` (indicating it can't handle the query), a remote API could be called for a more powerful model <sup>25</sup>. This indicates foresight in handling complex queries that a small model might not solve. The final step was adding **text-to-speech** output so that the generated answer could be heard by the user. For this, they used Microsoft's Edge TTS with a chosen voice (en-GB Sonia Neural) to produce a spoken version of the answer <sup>5</sup>. Testing this involved playing the audio and adjusting volume or pronunciation if needed.

Throughout development, the team tested the system in integrated fashion: saying the wake word, asking a full question, and observing the entire pipeline's behavior. Early integration tests might have revealed timing issues (for example, ensuring the camera frame capture happens after the question is asked, not before). The repository's `main.py` shows a loop that waits for the wake word, then orchestrates all components sequentially <sup>26</sup> <sup>27</sup>. Each integration step required debugging – e.g., if the LLM took too long or gave irrelevant answers, they would refine the prompt; if object detection misidentified something, they might filter by confidence or adjust lighting. One notable development decision visible in the code is filtering low-confidence detections (they only include objects with confidence  $\geq 0.6$  in the described content) <sup>28</sup>, which was likely added after noticing that false positives could confuse the user.

In summary, the development procedure was incremental: **(1)** get localization working, **(2)** add gesture recognition, **(3)** integrate object detection, **(4)** implement voice input/output, and **(5)** polish with language generation. Each step built on previous ones, turning separate experimental modules into a cohesive system. The result is a robust architecture where triggering one component (like the voice query) flows into activating all others, demonstrating careful integration of multiple machine learning and hardware interface technologies into one product.

## Function and Usage Introduction

GestureSpace is designed to be user-friendly and operates as an interactive assistant that one might use in daily life or specific scenarios. Here's how a typical usage session might unfold for a user:

**1. Activation via Voice:** The user begins by saying the wake word aloud, e.g., "Hey Gesture Space." The system's microphone is always listening passively for this phrase, and once detected <sup>26</sup>, the system becomes "active." This is analogous to saying "Hey Siri" or "OK Google" for those respective assistants, but with GestureSpace the wake word is custom and processed locally for privacy. Upon activation, the system gives a short audio cue or greeting (in the code, it is set to potentially play a prompt like "We are ready to assist you. Please showcase the environment around you." <sup>29</sup>). This lets the user know they can proceed with their question or command.

**2. Voice Query and Understanding:** The user then asks a question or states a request regarding their environment. For example, a visually impaired user might ask: "What objects are on the table in front of me?" or "Where is my water bottle?" GestureSpace records the speech, using a dynamic silence detector to determine when the user has finished speaking <sup>8</sup>. Immediately after, it transcribes the spoken words into text using the wav2vec2 speech recognition model. This transcription is done in a fraction of a second for a typical sentence. The system now has the user's query in text form (e.g., "what objects are on the table in

front of me”), which it can use to tailor its response. If the query is unclear or incomplete, the system’s language model can still attempt to infer intent from it, much like a smart assistant would.

**3. Visual Scene Capture:** As soon as the user finishes the question, the system captures an image of the scene through the camera feed <sup>30</sup>. In wearable use cases, this could be a camera on smart glasses or a phone camera held by the user. It’s important that the image represents what the user is currently facing or pointing toward. In practice, the user might be instructed (by that initial prompt or training) to look around or orient the camera toward the area of interest when asking the question. The system could also capture a short video clip for more robust analysis, but in the current design a single frame is used for speed.

**4. Object Detection and Identification:** Immediately after capturing the frame, GestureSpace runs the YOLO-based object detection on the image <sup>27</sup>. In real-time, this will highlight and label all recognizable objects. For instance, the system might identify “laptop, person, cup” if those are visible. Each object’s label and image position are recorded. At the same time, the system checks if any AprilTag markers are visible in the frame to anchor the scene’s coordinate space (especially useful if the environment is a known space like a home or office instrumented with tags on walls or furniture). If markers are found, it calculates the camera’s pose (position and angle) relative to them <sup>31</sup>. If no markers are present, the system can still operate, but spatial accuracy will be lower – it might then assume the user’s perspective as the origin without an absolute reference.

**5. Gesture and Focus Detection:** GestureSpace also processes the frame with MediaPipe to detect the user’s hand and index finger. If the user is pointing at something while asking, the system will pick up the direction of that point <sup>32</sup> <sup>33</sup>. For example, if the user points their finger toward a particular object (like a book on a shelf) while asking “What is this?”, the system can use the pointing direction to disambiguate which object the question refers to. In cases where multiple objects are detected, knowing the pointing direction helps filter out the irrelevant ones. The code gathers the direction vector of the finger and the 3D coordinates of where the finger is “aiming” on the reference plane <sup>33</sup>. In future versions, this could allow the system to say “You are pointing at [object].” Currently, the main loop does compute the pointing direction but doesn’t explicitly narrate it; it is an available signal that could improve answer relevance.

**6. Spatial Reasoning:** With object positions known, the system translates those into descriptions relative to the user. For each detected object, GestureSpace uses the camera pose to compute coordinates like “(x, y, z) = (1.2 m, -0.5 m, 0.0 m)” meaning for example an object is 1.2 meters in front (positive x), 0.5 meters to the left (negative y), and at ground level (z=0) <sup>14</sup>. Internally, it appends each object and its location into a text “content” string that will be fed to the language model. For example, this content might be: “cup at 0.50, 1.20, 0.00 (x, y, z). Confidence: 0.91\nlaptop at -0.30, 1.00, 0.00 (x, y, z). Confidence: 0.85\n”. This raw content is like a factual list of what’s around. Additionally, the user’s original question text is preserved. These together form the input to the next step.

**7. Natural Language Answer Generation:** GestureSpace now uses its built-in language model to generate an answer that directly addresses the user’s query, in a friendly manner. It combines the factual scene content with a pre-written prompt. The prompt ensures the model knows it should act as a descriptive assistant for someone who cannot see the whole room <sup>17</sup>. The user’s question is also included at the end of this input context. For instance, the composed prompt to the LLM might read: “*You are talking to a person who cannot see the room. You have a list of objects and their 3D coordinates. Describe the scene to them without using numbers, just saying if things are near or far. The list of objects: cup at (0.5,1.2,0.0), laptop at (-0.3,1.0,0.0). Now the user is asking: ‘What objects are on the table in front of me?’*”. Given this, the language model will

generate a response like: *“There is a laptop a little to your left on the table and a cup further to the right. The laptop is close by, and the cup is a bit farther away on the table.”* This step is where the system truly interprets the user’s request and the environmental data to produce a relevant answer. If the user had asked a different question (say, “Where is my water bottle?”), the model would specifically look for a “water bottle” in the content and respond accordingly, e.g., “Your water bottle is on the table, a few feet in front of you.” If the model for some reason cannot handle the query (which might happen for very complex instructions), the system is set up to optionally call a more powerful cloud AI (this is a backup mechanism not usually needed for straightforward descriptions).

**8. Auditory Feedback:** The final step is delivering the answer to the user through sound. GestureSpace’s text-to-speech engine converts the generated sentence into spoken words using a natural-sounding synthetic voice <sup>5</sup>. The choice of voice (British English “Sonia” voice, in this case) is meant to be clear and pleasant. The system plays this audio through the user’s device speakers or headphones. So, continuing the example, the user will hear something like: *“I see a laptop on the table right in front of you, and a cup a bit further to the right.”* This audible response is effectively the system acting as the user’s eyes, describing the scene. The entire loop from wake word to spoken answer happens quickly – potentially in just a few seconds – to feel like a real conversation. The user can then decide a next action, perhaps moving toward the mentioned object or asking another question. The system would go back to listening for the wake word again, ready for the next query.

**Usage Contexts:** GestureSpace is primarily envisioned for assistive use by people who are blind or have low vision. In practice, such a user could wear a camera (attached to glasses or a chest harness) and carry a small computing unit running this system. They might walk into a new room and say “Hey Gesture Space, what’s around me?” The device would then enumerate key objects and their positions, enabling the user to orient themselves. It could say, for instance, “There is a sofa about three steps ahead and a coffee table just in front of you.” This dramatically improves situational awareness compared to a traditional screen reader which can only read text on a device. Another scenario is in a smart home: the user could ask, “Gesture Space, where did I leave my keys?” and the system, recognizing a key-shaped object on the counter, could respond, “Your keys are on the kitchen counter to your left.” Additionally, the gesture component allows control and query by hand signals – for example, the user might make a certain gesture to prompt the system for an overview (if speaking is inconvenient) or point toward a shelf and ask “What’s on this shelf?” The AR integration means that if the user has some visual capability or uses an AR display, the system could highlight objects or overlay labels on their view. For a fully blind user, AR audio cues (like spatial sound coming from the direction of an object) could be used; GestureSpace could play the spoken description panned to the direction of the object relative to the user, giving an auditory sense of direction.

Finally, in low-visibility scenarios (not necessarily involving blindness, but situations like darkness or smoke), a user could similarly use GestureSpace. For example, a firefighter in a smoke-filled room might say, “What do you see?” and get a quick rundown: “Doorway ahead, furniture on the left,” aiding navigation when vision is impaired. In all cases, the usage of GestureSpace is meant to feel like interacting with a knowledgeable companion who can both see the environment and listen to the user’s needs, then guide them with voice feedback.

## Creative Points and Application Scenarios

GestureSpace introduces several creative innovations by fusing technologies that are typically separate. One key creative point is its **multi-modal interaction**: the system blends voice, vision, and gesture input in

a seamless way. Voice assistants alone (like Alexa or Siri) cannot perceive the physical environment, and computer vision systems alone do not take natural language input – GestureSpace bridges this gap. The user can speak naturally and use intuitive gestures, which is far more flexible and human-friendly than pressing buttons or using a keyboard. This design makes advanced technology accessible even to those with limited technical skills, since asking a question and pointing is something anyone can do.

Another innovative aspect is **spatially-aware scene description**. Unlike conventional object recognition apps that might just list objects, GestureSpace understands and communicates layout. It doesn't just say "there is a chair"; it says *where* the chair is (e.g. "to your left" or "near the window"). This is made possible by the integration of AprilTag-based mapping – effectively turning the environment into an indexed space. By planting a few inexpensive AprilTag markers in a room, the system gains a consistent frame of reference, enabling use cases like an indoor navigation aid (imagine stickers or magnets with AprilTags placed in key locations of a house or classroom – the device can then always tell the user, "You are in the entrance hall" or "The whiteboard is on the north wall"). This approach is like a mini indoor GPS and is quite novel in assistive tech, which historically hasn't combined fixed environment markers with AI in a user-facing tool.

The use of a **Large Language Model for tailored descriptions** is another creative point. Instead of hard-coding how the system should respond for every possible scenario, the developers let an AI model dynamically generate the description. This means the system can flexibly answer a wide range of questions about the environment, not just a single hardwired task. For example, if the user asks, "Is my coffee mug on the desk?", the system can specifically answer yes or no and give details, whereas if the user asks, "What does the room look like?", it can give a broad summary. The prompt design guiding the LLM to avoid numeric coordinates and use relative terms is particularly thoughtful <sup>17</sup> – it shows consideration of the end-user's needs (most people, especially if visually impaired, benefit from "left/right/near/far" more than "1.5 meters away"). In essence, GestureSpace can translate raw sensor data into a story-like explanation, which is cutting-edge in humanizing assistive technology.

**Application scenarios** for GestureSpace are diverse. The most immediate one is as an assistive device for the **blind or visually impaired**. Here, GestureSpace serves as an AI co-pilot for everyday life: helping users find objects, avoid obstacles, or just get a sense of their surroundings. It could be used in the home (finding items, identifying who just walked into the room) or in public spaces (e.g., a user could point their phone camera around and ask "Where's the exit?" in an unfamiliar building, and the system could respond based on recognizing exit signs or doors). Another scenario is in **Augmented Reality (AR) smart glasses** for sighted users: a traveler could wear AR glasses with GestureSpace and get information on landmarks by looking and pointing ("What is that building?" – the system detects the building and perhaps, if connected to an online database, names it and gives a brief history). This blends education with exploration. In a classroom or lab, **educational uses** are promising – students could use GestureSpace to interact with exhibits or experiments (point at a part of a science apparatus and ask what it is or what it does, receiving an explanation).

For **gesture-based computer control**, GestureSpace's technology could be adapted to let users with mobility impairments control devices. For instance, a user could perform a certain hand gesture to signal "next slide" during a presentation, and the system (with its gesture classifier) could execute that command on the computer. This is essentially carrying forward the gesture control research that was integrated into GestureSpace. By combining it with voice, even more control is possible: a user could either speak a command or gesture, depending on what is easiest at the moment (useful for someone who might not be able to speak or move at certain times, providing alternative inputs).



Another creative application is in **low-visibility and emergency scenarios**. Firefighters, as mentioned, or spelunkers in caves, could use a version of GestureSpace with infrared cameras. The system could identify hazards (e.g., “there is a staircase ahead” or “flames to your right”) and communicate them. In such a scenario, the wake word and voice interaction allow hands-free operation, which is critical when carrying equipment. The gesture part might be less used here, but the spatial awareness and descriptive ability are life-saving features.

**Social and service robotics** is yet another area: a robot equipped with GestureSpace’s vision-language system could interact more naturally with humans. For example, a home service robot could detect when a person points at something and listen to their instruction (“Bring me that”) – it would understand what “that” refers to by combining the pointing direction and object recognition. It could then respond via voice to confirm. This is a sophisticated capability that typical robots currently lack, and GestureSpace offers a template for it.

Moreover, the project’s approach expands the application range of accessibility tech. Traditional screen-readers focus on digital content accessibility, whereas GestureSpace extends accessibility to **physical environments**. A blind person with this tool can “read” a room much like they’d read a website with a screen-reader. This fusion of physical and digital accessibility is a frontier area, and GestureSpace is a creative step in that direction.

In conclusion, GestureSpace’s innovative fusion of gesture, voice, and visual scene understanding opens up a wide array of use cases. Whether it’s helping a visually impaired individual navigate their daily life, enhancing AR experiences for exploration and learning, or providing hands-free assistance in critical situations, the technology demonstrates how combining modalities can create a more intuitive and powerful user experience. Each component on its own (voice assistant, object detector, etc.) is useful, but together in GestureSpace, they form a synergistic system that behaves in a human-like interactive manner. This multi-modal synergy is the core creative leap of the project – it transforms how humans can interface with both machines and the world around them, making the environment itself more “readable” and responsive. <sup>13</sup> <sup>17</sup>

---

<sup>1</sup> <sup>2</sup> <sup>13</sup> <sup>15</sup> <sup>23</sup> [locate.py](#)

<https://github.com/7086cmd/GestureSpace/blob/1da575fadedb4b22a2076de2955df687706342db1/locate.py>

<sup>3</sup> <sup>4</sup> <sup>19</sup> <sup>20</sup> <sup>21</sup> [gesture.py](#)

<https://github.com/7086cmd/GestureSpace/blob/1da575fadedb4b22a2076de2955df687706342db1/gesture.py>

<sup>5</sup> [play.py](#)

<https://github.com/7086cmd/GestureSpace/blob/1da575fadedb4b22a2076de2955df687706342db1/play.py>

<sup>6</sup> [2024-11-15-prototype-combination-of-camera-pose-estimation-tracking-and-neural-networks.md](#)

<https://github.com/7086cmd/blog/blob/d4b77e412ae7eca017912e39dbadeb4b11f121a9/blog/2024-11-15-prototype-combination-of-camera-pose-estimation-tracking-and-neural-networks.md>

<sup>7</sup> [waking.py](#)

<https://github.com/7086cmd/GestureSpace/blob/1da575fadedb4b22a2076de2955df687706342db1/waking.py>

<sup>8</sup> <sup>9</sup> [recognition.py](#)

<https://github.com/7086cmd/GestureSpace/blob/1da575fadedb4b22a2076de2955df687706342db1/recognition.py>

10 18 **detection.py**

<https://github.com/7086cmd/GestureSpace/blob/1da575fadeb4b22a2076de2955df687706342db1/detection.py>

11 22 **objects.py**

<https://github.com/7086cmd/GestureSpace/blob/1da575fadeb4b22a2076de2955df687706342db1/objects.py>

12 32 33 **finger.py**

<https://github.com/7086cmd/GestureSpace/blob/1da575fadeb4b22a2076de2955df687706342db1/finger.py>

14 26 27 28 29 30 31 **main.py**

<https://github.com/7086cmd/GestureSpace/blob/1da575fadeb4b22a2076de2955df687706342db1/main.py>

16 24 25 **local.py**

<https://github.com/7086cmd/GestureSpace/blob/1da575fadeb4b22a2076de2955df687706342db1/local.py>

17 **shared.py**

<https://github.com/7086cmd/GestureSpace/blob/1da575fadeb4b22a2076de2955df687706342db1/shared.py>